

# OpenStack Havana

Joseph Callen

December 12, 2013

### **Acknowledgements**

I have read many articles from various sources I am sure I have missed to source them some along the way.

# Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
1.1	Physical Hosts . . . . .	2
1.2	oVirt . . . . .	2
1.2.1	oVirt Engine Install . . . . .	3
1.2.2	PXE . . . . .	3
1.2.3	DHCP . . . . .	4
1.2.4	oVirt Node Install . . . . .	5
<b>2</b>	<b>All-In-One install with Neutron using VLANs</b>	<b>8</b>
2.1	Physical Networking . . . . .	8
2.2	Installation . . . . .	10
2.2.1	Packstack . . . . .	10
2.3	Neutron Networking . . . . .	10
2.3.1	External . . . . .	10
2.3.2	Tenant . . . . .	11
2.3.3	Troubleshooting . . . . .	11
<b>3</b>	<b>Multinode install using Neutron with GRE</b>	<b>20</b>
3.1	Physical Networking . . . . .	20
3.2	Installation . . . . .	21
3.2.1	DNS . . . . .	21
3.2.2	Virtual Machine Deployment . . . . .	21
3.2.3	SSH Keys . . . . .	22
3.2.4	Packstack . . . . .	22
3.3	Neutron Networking . . . . .	24
3.3.1	External . . . . .	26
3.3.2	Tenant . . . . .	26
3.3.3	Troubleshooting . . . . .	27
3.4	OpenDaylight . . . . .	27
3.4.1	Install Controller . . . . .	28

3.4.2	Configure Open vSwitch . . . . .	29
<b>4</b>	<b>Instances</b>	<b>30</b>
4.1	Security Group Configuration . . . . .	30
4.2	Create Images . . . . .	31
4.2.1	Ubuntu LTS 12.04 . . . . .	31
4.2.2	Windows 2012 R2 . . . . .	31
4.3	Create Instance . . . . .	32
<b>A</b>	<b>Notes</b>	<b>33</b>
A.1	Neutron . . . . .	33
A.2	Horizon . . . . .	33
A.3	Cinder . . . . .	33
A.4	Nova . . . . .	34
<b>B</b>	<b>Links</b>	<b>35</b>

# List of Figures

1.1	System Volumes . . . . .	5
1.2	Network Configuration . . . . .	5
1.3	IP Configuration . . . . .	6
1.4	oVirt Engine Configuration . . . . .	6
1.5	Hosts in oVirt Engine . . . . .	7
2.1	Compute Node . . . . .	9
2.2	Cisco Catayst 4006 . . . . .	9
2.3	VLAN Scenario . . . . .	12
3.1	OpenStack Node Details . . . . .	24
3.2	OpenDaylight Controller UI . . . . .	28

# List of Tables

1.1 Physical Compute . . . . .	2
--------------------------------	---

# Listings

1.1	Install oVirt repo and package . . . . .	3
1.2	Download and convert oVirt Node . . . . .	3
1.3	Copy tftproot to root, restore SELinux context, change permissions . . . . .	3
1.4	Enable TFTP . . . . .	3
1.5	Install ISC DHCP server . . . . .	4
1.6	Example DHCPd configuration . . . . .	4
1.7	Restart DHCPd . . . . .	4
1.8	Disable SELinux . . . . .	6
2.1	VLAN Configuration . . . . .	8
2.2	Instance uplink interface; gi2/5; #1 . . . . .	8
2.3	Management interface; gi3/13; #2 . . . . .	8
2.4	Configure Packstack . . . . .	10
2.5	Answer file changes . . . . .	10
2.6	Execute Packstack . . . . .	10
2.7	Create network . . . . .	10
2.8	Create subnet . . . . .	11
2.9	Create router . . . . .	11
2.10	Set gateway . . . . .	11
2.11	Create network . . . . .	11
2.12	Create subnet . . . . .	11
2.13	Add interface . . . . .	11
2.14	map_router.py script . . . . .	11
2.15	map_router output . . . . .	18
3.1	Nova Node uplink interface; gi2/5 . . . . .	20
3.2	Virtual Machine uplink; gi3/4 . . . . .	20
3.3	Virtual Machine uplink; gi3/20 . . . . .	20
3.4	Remove OpenStack from all nodes . . . . .	21
3.5	DNS Entries . . . . .	21
3.6	Upload ISO . . . . .	21
3.7	Disable firewalld, install iptables . . . . .	22
3.8	Disable SELinux . . . . .	22

3.9	SSH Key Copy . . . . .	22
3.10	Packstack remote install after big hammer . . . . .	22
3.11	Configure Packstack . . . . .	23
3.12	Changes in Packstack answer file . . . . .	23
3.13	Execute Packstack . . . . .	23
3.14	Nova Node Open vSwitch . . . . .	24
3.15	Neutron Node Open vSwitch . . . . .	25
3.16	Create public network and subnet . . . . .	26
3.17	Create admin tenant router, network and subnet . . . . .	26
3.18	Set the router gateway . . . . .	26
3.19	Stop Neutron Services . . . . .	27
3.20	Find the MariaDB root password . . . . .	27
3.21	Login to MariaDB . . . . .	27
3.22	Get the ID for the floating IP and delete . . . . .	27
3.23	Install Oracle Java JDK . . . . .	28
3.24	Disable firewalld . . . . .	28
3.25	Get the latest controller . . . . .	28
3.26	Neutron node set controller and manager . . . . .	29
3.27	Nova node set controller . . . . .	29
4.1	Delete existing security group rules . . . . .	30
4.2	Delete existing security groups . . . . .	30
4.3	Windows 2012 R2 Image Create . . . . .	31
4.4	Add public key . . . . .	32
4.5	Windows administrator password . . . . .	32
A.1	Change Cinder volume clear . . . . .	33



# Executive Summary

OpenStack can be difficult software package to install and configure for a novice. I have created this document to guide someone through the install that I performed in the lab. Hopefully with a "real world" example it should be easier to understand the various modules and concepts required for an OpenStack deployment. This document was created with L<sup>A</sup>T<sub>E</sub>X.

# Chapter 1

## Introduction

### 1.1 Physical Hosts

In order to test the installation and operational aspects of a true OpenStack environment I needed additional nodes. Utilizing old laptop hardware I was able to create four additional virtual machines to run OpenStack services on. The compute node remains physical. I am still using the RDO Havana community edition on Fedora 19 x64.

Table 1.1: Physical Compute

Hardware		Dev	Model	Network	Usage
AMD Athlon Dual Core	1	eth0	r8169	Managment	Nova, Glance &
8 GB RAM, 2 TB HD	2	eth1	e1000	Instance uplink	Cinder
Levono T61	1	eth0	e1000e	Management &	oVirt Node
Intel®Core2™Duo CPU T7100, 4 GB RAM				Instance Uplink	
Dell Latitude D620	1	eth0	tg3	Management &	oVirt Node
Intel®Core2™Duo CPU T5500, 4 GB RAM				Instance Uplink	
Dell PE2850	1	em1	e1000	Management	NFS
Intel®Xeon™, 2 GB RAM, 3x146GB 10K				Instance Uplink	

### 1.2 oVirt

oVirt is a management and host node virtualization package similar to vSphere. It uses Linux KVM as the hypervisor with either Fedora or CentOS as the base OS. The management console is web based.

### 1.2.1 oVirt Engine Install

Before installation DNS needs to be configured, including an A record of the oVirt engine and nodes.

Listing 1.1: Install oVirt repo and package

---

```
1 yum localinstall http://ovirt.org/releases/ovirt-release-fedora.noarch.  
   rpm  
2 yum install -y ovirt-engine  
3 engine-setup
```

---

### 1.2.2 PXE

To install an oVirt node the easiest method I have found is use PXE. First lets convert the ISO to tftpboot. **Note:** This is an older release of ovirt-node, the current version has a bug with installation.<sup>1</sup>

Listing 1.2: Download and convert oVirt Node

---

```
1 wget http://resources.ovirt.org/releases/3.3/iso/ovirt-node-iso  
   -3.0.1-1.0.2.vdsm.fc19.iso  
2 livedvd-iso-to-pxeboot ovirt-node-iso-3.0.1-1.0.2.vdsm.fc19.iso
```

---

Listing 1.3: Copy tftpboot to root, restore SELinux context, change permissions

---

```
1 cp -R /tftpboot/ /  
2 restorecon -Rv /tftpboot/  
3 chmod -R 755 /tftpboot/
```

---

Listing 1.4: Enable TFTP

---

```
1 sed -i 's/disable.*=\ yes/disable\t\t\t= no/' /etc/xinetd.d/tftp  
2 # or  
3 vi /etc/xinetd.d/tftp  
4 service tftp  
5 {  
6     ...  
7     disable = no #change from yes to no  
8     ...  
9 }  
10 systemctl restart xinetd
```

---

---

<sup>1</sup>Red Hat Bugzilla Bug #1032228

### 1.2.3 DHCP

The DHCP server needs to be configured to provide addresses and PXE related configuration elements.

Listing 1.5: Install ISC DHCP server

---

```
1 yum install dhcp -y
```

---

Below is a partial example of my DHCPd configuration file. The most important section is "ovirt-server1". Options filename and next-server must be listed for PXE boot to function correctly. The option filename will not change and should be listed as "pxelinux.0".

Listing 1.6: Example DHCPd configuration

---

```
1 authoritative;
2 ddns-update-style none;
3
4 ddns-domainname      "virtomation.com";
5 option domain-name   "virtomation.com";
6 option domain-name-servers 10.53.252.123, 10.53.252.246;
7 option ntp-servers   10.53.252.123;
8 default-lease-time  86400;
9 option ip-forwarding off;
10 ...
11 subnet 10.53.253.0 netmask 255.255.255.0 {
12     range 10.53.253.130 10.53.253.140;
13     option routers 10.53.253.1;
14     option subnet-mask 255.255.255.0;
15 }
16
17 # HOST - RHEV
18 host ovirt-server1 {
19     hardware ethernet 00:15:C5:59:AE:36; # MAC from machine that
20         will boot via PXE
21     fixed-address 10.53.253.50; # Fixed IP address
22     filename "pxelinux.0";
23     next-server 10.53.252.50; # DHCP Server
24     option host-name "virkvmpaw001.virtomation.com";
25 }
```

---

After the configuration has been saved remember to restart dhcpd.

Listing 1.7: Restart DHCPd

---

```
1 systemctl reload-or-try-restart dhcpd
```

---

### 1.2.4 oVirt Node Install

The installation of the oVirt node is generally as easy as ESXi. I am only displaying figures that need extra explanation. Figure 1.1 shows the system volumes and sizes. If you are using a USB flash drive set the log and swap to 5 MB.

Figure 1.1: System Volumes

```

oVirt Node Hypervisor 3.0.3-1.1.fc19

Storage Volumes

Please enter the sizes for the following partitions in MB

UEFI/Bios:                256_____
Root & RootBackup:        512_____
(2 partitions at 512MB each)

Swap:                      5_____
Config:                    5_____
Logging:                   5_____
Data:                      -1_____

< Quit >    < Back >    < Next >

```

The interfaces are not automatically configured. Figure 1.2 shows the available interfaces, select an interface and hit enter. In the next screen listed in Figure 1.3 configure the IP settings of the selected interface.

Figure 1.2: Network Configuration

```

oVirt Node Hypervisor 3.0.3-1.1.fc19

Status
Network
Security
Keyboard
Logging
Kdump
Remote Storage
Monitoring
Diagnostics
oVirt Engine
Performance
Plugins

System Identification

Hostname:                localhost_____
DNS Server 1:             _____
DNS Server 2:             _____
NTP Server 1:             0.fedora.pool.ntp.org_____
NTP Server 2:             1.fedora.pool.ntp.org_____

Available System NICs
Device  Status  Model  MAC Address
eth0    Unconfigured Red Hat, Inc 52:54:00:a5:94:01

(1 / 1)

< Ping >    < Create Bond >
< Save >    < Reset >

```

Finally once the oVirt is available on the network you can configure it with the oVirt engine. As displayed in Figure 1.4 configure the management server using either FQDN or IP and retrieve the certificate. You can also set the optional password, which I recommend

Figure 1.3: IP Configuration

```

NIC Details: eth0
Driver:      virtio_net
Link Status: Disconnected
Vendor:      Red Hat, Inc
MAC Address: 52:54:00:a5:94:01

IPv4 Settings
Bootprotocol: (X) Disabled ( ) DHCP ( ) Static
IP Address:   _____ Netmask:   _____
Gateway:     _____

IPv6 Settings
Bootprotocol: (X) Disabled ( ) Auto ( ) DHCP ( ) Static
IP Address:   _____ Prefix Length: _____
Gateway:     _____

VLAN ID:     _____

Use Bridge:   [ ]

< Flash Lights to Identify >
< Save >      < Close >

```

for a lab environment since it will also start ssh and set the root password.

Figure 1.4: oVirt Engine Configuration

```

oVirt Node Hypervisor 3.0.3-1.1.fc19

Status
Network
Security
Keyboard
Logging
Kdump
Remote Storage
Monitoring
Diagnostics
oVirt Engine
Performance
Plugins

oVirt Engine Configuration

Management Server:
Management Server Port: 443

< Retrieve Certificate >
Certificate Status: N/A

Optional password for adding Node through oVirt Engine UI
Note: Setting password will enable SSH daemon
Password:
Confirm Password:

< Save & Register >

```

**Note:** In my testing SELinux needed to be disabled before for any virtual machines would start. To disable SELinux use the command listed below. Unfortunately I haven't found a way to persist the changes after a reboot. I would guess that re-mounting the filesystem read-write and changing setting in `/etc/selinux/config` file would work but I have yet to try it.

Listing 1.8: Disable SELinux

---

```
1 setenforce 0
```

---

Below is a screenshot of the oVirt Administration UI with the available hosts listed.

Figure 1.5: Hosts in oVirt Engine



## Chapter 2

# All-In-One install with Neutron using VLANs

### 2.1 Physical Networking

I have three VLAN 98 — 100 configured on a dot1q trunk to interface p9p1 on a Fedora 19 host. VLAN 100 has a gateway IP configured (10.53.100.1) on the Cisco Catalyst 4006. VLAN 98 and 99 are non-routed vlans, at least as configured on the catalyst, so they have no gateway IP addresses configured, that will be provided by the neutron router.

---

Listing 2.1: VLAN Configuration

---

```
1 interface Vlan98
2   no ip address
3 end
4 interface Vlan99
5   no ip address
6 end
7 interface Vlan100
8   ip address 10.53.100.1 255.255.255.0
9   ip helper-address 10.53.252.50
10 end
```

---

---

Listing 2.2: Instance uplink interface; gi2/5; #1

---

```
1 interface GigabitEthernet2/5
2   switchport trunk encapsulation dot1q
3   switchport trunk allowed vlan 2-4,12,98-100,250-253
4   switchport mode trunk
5 end
```

---

---

Listing 2.3: Management interface; gi3/13; #2

---



Figure 2.1: Compute Node

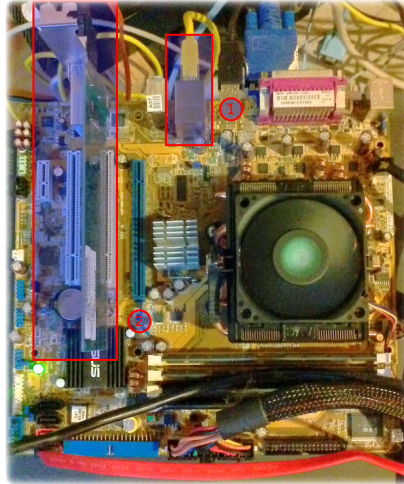
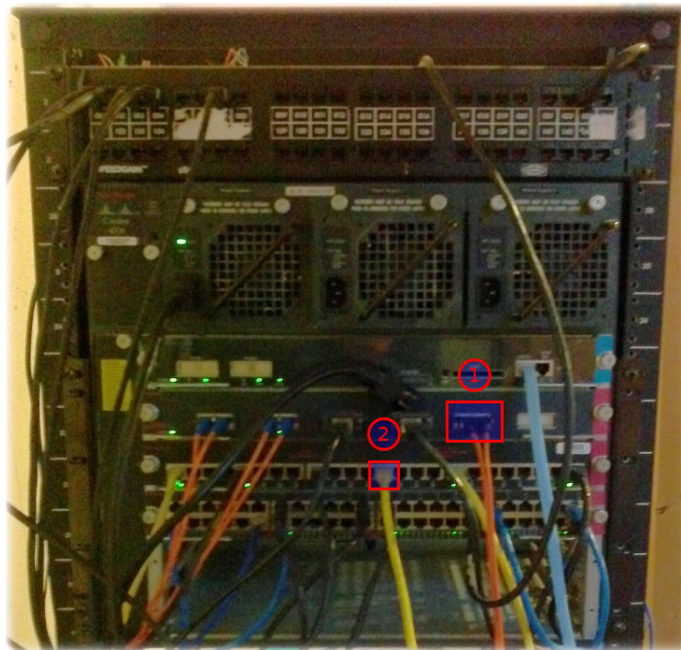


Figure 2.2: Cisco Catalyst 4006



```

1 WS-C4006-SUPV#sh run int gi3/13
2 Building configuration...
3
4 Current configuration : 65 bytes
5 !
6 interface GigabitEthernet3/13
7   switchport access vlan 253
8 end

```

---

## 2.2 Installation

### 2.2.1 Packstack

Listing 2.4: Configure Packstack

```

1 DATE=`date +"%Y_%m_%d_%H_%M_%S"`
2 packstack --gen-answer-file=packstack-${DATE}
3 vi packstack-${DATE}

```

---

Listing 2.5: Answer file changes

```

1 CONFIG_NEUTRON_INSTALL=y
2 CONFIG_CINDER_VOLUMES_CREATE=n
3 CONFIG_NEUTRON_L3_EXT_BRIDGE=provider
4 CONFIG_NEUTRON_OVS_TENANT_NETWORK_TYPE=vlan
5 CONFIG_NEUTRON_OVS_VLAN_RANGES=inter-vlan:98:99
6 CONFIG_NEUTRON_OVS_BRIDGE_MAPPINGS=inter-vlan:br-inst
7 CONFIG_NEUTRON_OVS_BRIDGE_IFACES=br-inst:p9p1

```

---

Listing 2.6: Execute Packstack

```

1 packstack --answer-file=packstack-${DATE}

```

---

**Note:** Monitor the progress in detail via `/var/tmp/packstack/`

## 2.3 Neutron Networking

### 2.3.1 External

Now we will configure VLAN 100 to be our external network. First we need to source the keystone file created by packstack.

Listing 2.7: Create network

```

1 neutron net-create ext_vlan100 --provider:network_type vlan --provider:
   physical_network inter-vlan --provider:segmentation_id 100 --router:
   external=True

```

---

Create a subnet for ext\_vlan100, this will be used for floating addresses.

Listing 2.8: Create subnet

---

```
1 neutron subnet-create ext_vlan100 --gateway 10.53.100.1 --name
   ext_vlan100_subnet 10.53.100.0/24 --allocation-pool start
   =10.53.100.10,end=10.53.100.253 --enable_dhcp=False
```

---

Listing 2.9: Create router

---

```
1 neutron router-create ext_vlan100_router
```

---

Listing 2.10: Set gateway

---

```
1 neutron router-gateway-set ext_vlan100_router ext_vlan100
```

---

### 2.3.2 Tenant

Listing 2.11: Create network

---

```
1 TENANT_ID=$(keystone tenant-list|awk '/admin/ {print $2}')
2 neutron net-create --tenant-id ${TENANT_ID} net01
```

---

Listing 2.12: Create subnet

---

```
1 neutron subnet-create net01 --name net01_subnet 192.168.98.0/24
```

---

Listing 2.13: Add interface

---

```
1 neutron router-interface-add ext_vlan100_router net01_subnet
```

---

The ext\_vlan\_router then becomes the default gateway for the vlan with an IP of 192.168.98.1. And the result as displayed in Horizon.

### 2.3.3 Troubleshooting

I believe OpenStack networking with Neutron is the most difficult part of the stack to understand. The problem being just to have isolated tenant networks requires linux bridge, open vswitch bridge, and virtual Ethernet devices. I have created a python script to better understand the relationship between a virtual machine and those devices.

The table below shows the interfaces between the vm and int-br-eth (as shown in Figure 2.3), where a majority of the devices are configured.

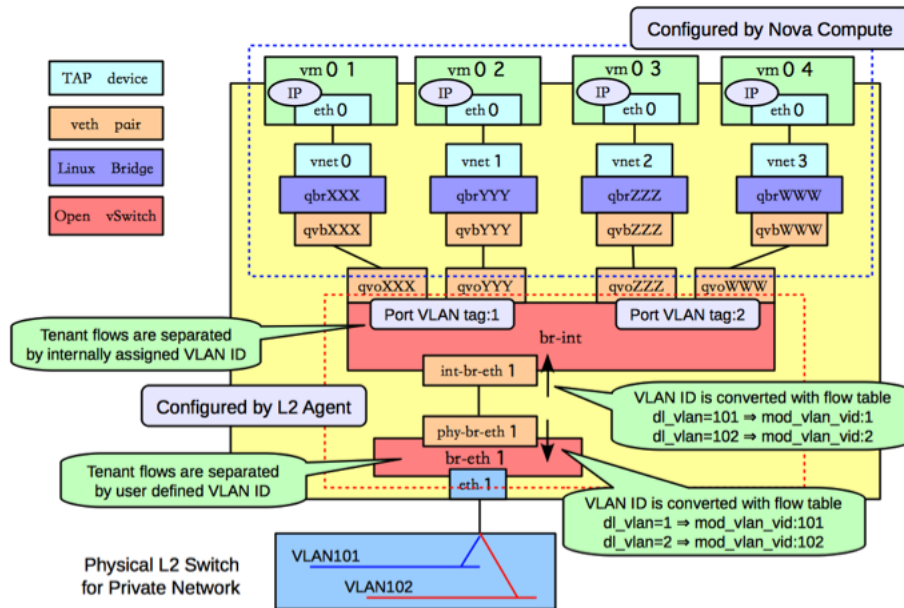
Listing 2.14: map\_router.py script

---

```
1 #Sources:
2 # http://effbot.org/zone/element-xpath.htm
```

---

Figure 2.3: VLAN Scenario



```

3 # http://eli.thegreenplace.net/2012/03/15/processing-xml-in-python-with
  -elementtree/
4 # http://wiki.libvirt.org/page/SSHSetup
5 # http://docs.python.org/2/tutorial/datastructures.html#dictionaries
6 # http://libvirt.org/git/?p=libvirt.git;a=blob;f=tools/virsh-domain-
  monitor.c
7 # http://www.linuxproblem.org/art_9.html
8 # http://j2labs.tumblr.com/post/4477180133/ssh-with-pythons-paramiko
9
10 __author__ = 'jcallen'
11 import os
12 import libvirt
13 import sys
14 import prettytable
15 import novaclient.v1_1.client as nvclient
16 from neutronclient.v2_0 import client
17 import keystoneclient.v2_0.client as ksclient
18 import xml.etree.ElementTree as ET
19 import paramiko
20 import re
21 import argparse
22
23
24 def keystone_connect(hostname, username, password, tenant):
25     auth_url = "http://%s:35357/v2.0" % hostname

```

```
26
27     keystone = ksclient.Client(auth_url=auth_url,
28                                username=username,
29                                password=password,
30                                tenant_name=tenant)
31     if keystone is None:
32         print 'Failed to open connection to OpenStack instance'
33         sys.exit(1)
34
35     return keystone
36
37
38 def neutron_connect(hostname, keystone):
39     os_url = "http://%s:9696/" % hostname
40     neutron_conn = client.Client(endpoint_url=os_url, token=keystone.
41                                  auth_token)
42     if neutron_conn is None:
43         print 'Failed to open connection to OpenStack instance'
44         sys.exit(1)
45
46     return neutron_conn
47
48 def nova_connect(hostname, username, password, tenant):
49     authurl = "http://%s:5000/v2.0" % hostname
50     nova_conn = nvclient.Client(username, password, tenant, authurl,
51                                 service_type="compute")
52     if nova_conn is None:
53         print 'Failed to open connection to OpenStack instance'
54         sys.exit(1)
55
56     return nova_conn
57
58 def qemu_connect(hostname, username):
59     qemu = "qemu+ssh://%s@%s/system" % (username, hostname)
60     qemu_conn = libvirt.openReadOnly(qemu)
61     if qemu_conn is None:
62         print 'Failed to open connection libvirt'
63         sys.exit(1)
64
65     return qemu_conn
66
67
68 def ssh_connect(host, username, private_key, port=22):
69     """Helper function to initiate an ssh connection to a host."""
70     transport = paramiko.Transport((host, port))
71
72     if os.path.exists(private_key):
73         rsa_key = paramiko.RSAKey.from_private_key_file(private_key)
```

```

74         transport.connect(username=username, pkey=rsa_key)
75     else:
76         raise TypeError("Incorrect private key path")
77
78     return transport
79
80
81 def exec_cmd(transport, command):
82     """Executes a command on the same server as the provided
83     transport
84     """
85     try:
86         channel = transport.open_session()
87         channel.exec_command(command)
88         if channel.recv_exit_status() == 0:
89             output = channel.makefile('rb', -1).readlines()
90             return output
91         else:
92             stderr_output = channel.makefile_stderr('rb', -1).readlines()
93             ()
94             print stderr_output
95             print "error: " + channel.recv_exit_status()
96             return ""
97     except:
98         print sys.exc_info()
99
100 def show_brctl_veth(device, transport):
101     BRCTL = "/usr/sbin/brctl"
102     command = BRCTL + " show " + device
103     output = exec_cmd(transport, command)
104     if output != "":
105         ifline = output[1]
106         match = re.findall("(qvb[\\w-]+$)", ifline)
107         veth = match[0]
108         return veth
109
110
111 def ethtool_adapter_stats(device, transport):
112     ETHTOOL = "/usr/sbin/ethtool"
113     command = ETHTOOL + " -S " + device
114     output = exec_cmd(transport, command)
115     if output != "":
116         ifline = output[1]
117         match = re.findall("peer_ifindex: ([\\d]+)$", ifline)
118         peer_ifindex = match[0]
119         return peer_ifindex
120
121
122 def ip_link(peer_ifindex, transport):

```

```

123     ETHTOOL = "/usr/sbin/ip"
124     command = ETHTOOL + " link "
125     output = exec_cmd(transport, command)
126     if output != "":
127         for interfaces in output:
128             ifline = interfaces
129             matchstring = "^%s:\s([\w-]+)" % peer_ifindex
130             match = re.findall(matchstring, ifline)
131             if len(match) != 0:
132                 veth_pair2 = match[0]
133                 return veth_pair2
134
135
136 def ovs_ofctl(action, device, transport):
137     """
138     Executes ovs-ofctl via paramiko ssh client
139
140     Requires a sudoers entry:
141     Defaults !requiretty
142     user ALL = NOPASSWD: /usr/bin/ovs-ofctl
143     """
144     OFCTL="sudo /usr/bin/ovs-ofctl"
145     command = OFCTL + " " + action + " " + device
146     output = exec_cmd(transport, command)
147     return output
148
149
150 def ovs_ofctl_dump_flows(device, tag, transport):
151     output = ovs_ofctl("dump-flows", device, transport)
152     if output != "":
153         for line in output:
154             matchstring = "mod_vlan_vid:%s" % tag
155             if re.search(matchstring, line) is not None:
156                 matchstring = "dl_vlan=(\d+)"
157                 match = re.findall(matchstring, line)
158                 if len(match) != 0:
159                     #print match
160                     vlan = match[0]
161                     return vlan
162
163
164 def ovs_vsctl(action, device, transport):
165     """
166     Executes ovs-vsctl via paramiko ssh client
167
168     Requires a sudoers entry:
169     Defaults !requiretty
170     user ALL = NOPASSWD: /usr/bin/ovs-vsctl
171     """
172

```

```

173     VSCTL="sudo /usr/bin/ovs-vsctl"
174     command = VSCTL + " " + action + " " + device
175     output = exec_cmd(transport, command)
176     return output
177
178
179 def ovs_vsctl_list_port(device, transport):
180     output = ovs_vsctl("list port", device, transport)
181     if output != "":
182         for line in output:
183             matchstring = "^tag[\s:]* (\d+)"
184             match = re.findall(matchstring, line)
185             if len(match) != 0:
186                 tag = match[0]
187                 return tag
188
189
190 def ovs_vsctl_port_to_br(device, transport):
191     output = ovs_vsctl("port-to-br", device, transport)
192     bridge = output[0].split("\n")[0]
193     return bridge
194
195
196 def get_router_id(neutron_conn, net_id):
197     """
198     Yes, I know there is a bug here, will resolve later!
199     """
200     net = neutron_conn.show_network(net_id)
201
202     for subnet in net['network']['subnets']:
203         ports = neutron_conn.list_ports(retrieve_all=True, subnet_id=
            subnet)
204         for port in ports['ports']:
205             if port['device_owner'] == 'network:router_interface':
206                 return port['device_id']
207
208
209 def domiflist(domxml):
210     tree = ET.ElementTree(ET.fromstring(domxml))
211     elements = tree.iterfind("./devices/interface")
212
213     for ele in elements:
214         interface_type = ele.attrib['type']
215         if interface_type == "bridge":
216             bridge = ele.find("source[@bridge]").attrib["bridge"]
217             tap = ele.find("target[@dev]").attrib["dev"]
218             mac = ele.find("mac[@address]").attrib["address"]
219             return bridge, tap, mac
220         else:
221             print "Error: Currently will only work with bridged

```



```

222         interfaces."
223
224 def args():
225     parser = argparse.ArgumentParser()
226     parser.add_argument('--hostname', help='Hostname of OpenStack
227         instance', required=True)
228     parser.add_argument('--os-username', help='OpenStack Username',
229         required=True)
230     parser.add_argument('--os-password', help='OpenStack Password',
231         required=True)
232     parser.add_argument('--os-tenant-name', help='OpenStack Tenant',
233         required=True)
234     parser.add_argument('--username', help='Username of OpenStack host',
235         required=True)
236     parser.add_argument('--priv-key', help='Location of private SSH Key',
237         required=True)
238     arguments = parser.parse_args()
239     return arguments
240
241 def main():
242     args_namespace = args()
243
244     nova_conn = nova_connect(args_namespace.hostname,
245         args_namespace.os_username,
246         args_namespace.os_password,
247         args_namespace.os_tenant_name)
248     qemu_conn = qemu_connect(args_namespace.hostname, args_namespace.
249         username)
250     keystone_conn = keystone_connect(args_namespace.hostname,
251         args_namespace.os_username,
252         args_namespace.os_password,
253         args_namespace.os_tenant_name)
254     neutron_conn = neutron_connect(args_namespace.hostname,
255         keystone_conn)
256
257     try:
258         pt = prettytable.PrettyTable(["Property", "Value"])
259         pt.align = "l"
260         servers = nova_conn.servers.list(detailed=True)
261         for srv in servers:
262             interface_list = srv.interface_list()
263             for interface in interface_list:
264                 router_id = get_router_id(neutron_conn, interface.
265                     net_id)
266                 print ("\nTroubleshooting commands:")
267                 print ("ip netns exec qrouter-%s ip a" % router_id)
268                 print ("ip netns exec qrouter-%s ip r" % router_id)

```

```

262         print ("ip netns exec qrouter-%s iptables -t nat -L -nv
                " % router_id)
263         print ("ip netns exec qrouter-%s iptables -S -t nat" %
                router_id)
264         print ("ip netns exec qrouter-%s ping 8.8.8.8" %
                router_id)
265
266         dom = qemu_conn.lookupByUUIDString(srv.id) # get domain
                from UUID
267         xml = dom.XMLDesc()
268         bridge, tap, mac = dom.iflist(xml)
269
270         transport = ssh_connect(args_namespace.hostname,
                args_namespace.username, args_namespace.priv_key)
271         veth_pair1 = show_brctl_veth(bridge, transport)
272         peer_ifindex = ethtool_adapter_stats(veth_pair1, transport)
273         veth_pair2 = ip_link(peer_ifindex, transport)
274         ovsbridge = ovs_vsctl_port_to_br(veth_pair2, transport)
275         tag = ovs_vsctl_list_port(veth_pair2, transport)
276         vlan = ovs_ofctl_dump_flows(ovsbridge, tag, transport)
277
278         pt.add_row(["OpenStack Name", srv.human_id])
279         pt.add_row(["QEMU Name", dom.name()])
280         pt.add_row(["UUID", dom.UUIDString()])
281         pt.add_row(["tap", tap])
282         pt.add_row(["linuxbridge", bridge])
283         pt.add_row(["MAC Address", mac])
284         pt.add_row(["VETH Pair #1", veth_pair1])
285         pt.add_row(["VETH Pair #2", veth_pair2])
286         pt.add_row(["ovsbridge", ovsbridge])
287         pt.add_row(["TAG", tag])
288         pt.add_row(["VLAN", vlan])
289
290         print(pt)
291         pt.clear_rows()
292
293     except:
294         print sys.exc_info()
295         sys.exit(1)
296
297 if __name__ == '__main__':
298     main()

```

Listing 2.15: map-router output

---

```

1 Troubleshooting commands:
2 ip netns exec qrouter-2672c623-3b4e-4c28-9f4f-505b17a25542 ip a
3 ip netns exec qrouter-2672c623-3b4e-4c28-9f4f-505b17a25542 ip r
4 ip netns exec qrouter-2672c623-3b4e-4c28-9f4f-505b17a25542 iptables -t
  nat -L -nv

```

```

5 ip netns exec qrouter-2672c623-3b4e-4c28-9f4f-505b17a25542 iptables -S
   -t nat
6 ip netns exec qrouter-2672c623-3b4e-4c28-9f4f-505b17a25542 ping 8.8.8.8
7 +-----+
8 | Property          | Value                               |
9 +-----+
10 | OpenStack Name    | ubuntu1204                         |
11 | QEMU Name         | instance-00000016                 |
12 | UUID              | e2c0d2b3-05c6-45e9-ad52-6f11647badbb |
13 | tap                | tap0598ef4e-cd                    |
14 | linuxbridge       | qbr0598ef4e-cd                    |
15 | MAC Address       | fa:16:3e:65:a7:1f                 |
16 | VETH Pair #1      | qvb0598ef4e-cd                    |
17 | VETH Pair #2      | qvo0598ef4e-cd                    |
18 | ovsbridge         | br-int                             |
19 | TAG               | 2                                  |
20 | VLAN              | 98                                  |
21 +-----+

```

## Chapter 3

# Multinode install using Neutron with GRE

### 3.1 Physical Networking

The configuration of the physical network is significantly reduced when using tunneling. To make configuration easier I only used VLAN 253.

Listing 3.1: Nova Node uplink interface; gi2/5

---

```
1 interface GigabitEthernet2/5
2   switchport access vlan 253
3   switchport mode access
4   !
```

---

Listing 3.2: Virtual Machine uplink; gi3/4

---

```
1 interface GigabitEthernet3/4
2   switchport access vlan 253
3   switchport mode access
4   !
```

---

Listing 3.3: Virtual Machine uplink; gi3/20

---

```
1 interface GigabitEthernet3/20
2   switchport access vlan 253
3   switchport mode access
4   !
```

---

## 3.2 Installation

First things first, we need to completely remove the OpenStack install on the physical compute node. To accomplish that I used the big hammer method.<sup>1</sup> This would also be useful if Packstack has issues with installation and you would like to start over without removing a snapshot or redeploying.

Listing 3.4: Remove OpenStack from all nodes

---

```

1 ssh root@virospaw002 'bash -s' < big_hammer.sh
2 for i in {1..3}; do ssh root@virospaw01$i 'bash -s' < big_hammer.sh;
   done

```

---

### 3.2.1 DNS

My current DNS servers are Windows 2012. To add multiple entries its just easier to use PowerShell, which is listed below.

Listing 3.5: DNS Entries

---

```

1 Add-DnsServerResourceRecordA -Name "virospaw002" -ZoneName "
   virtomation.com" -IPv4Address 10.53.253.70 -ComputerName
   10.53.252.123
2 Add-DnsServerResourceRecordA -Name "virospaw010" -ZoneName "
   virtomation.com" -IPv4Address 10.53.253.90 -ComputerName
   10.53.252.123
3 Add-DnsServerResourceRecordA -Name "virospaw011" -ZoneName "
   virtomation.com" -IPv4Address 10.53.253.100 -ComputerName
   10.53.252.123
4 Add-DnsServerResourceRecordA -Name "virospaw012" -ZoneName "
   virtomation.com" -IPv4Address 10.53.253.110 -ComputerName
   10.53.252.123
5 Add-DnsServerResourceRecordA -Name "virospaw013" -ZoneName "
   virtomation.com" -IPv4Address 10.53.253.120 -ComputerName
   10.53.252.123

```

---

### 3.2.2 Virtual Machine Deployment

Just like with vSphere we will create a template to use for virtual machine deployment. First we need to upload the Fedora 19 x64 iso to the NFS domain.

Listing 3.6: Upload ISO

---

```

1 wget http://download.fedoraproject.org/pub/fedora/linux/releases/19/
   Fedora/x86_64/iso/Fedora-19-x86_64-netinst.iso
2 ovirt-iso-uploader -i ISO_NFS upload Fedora-19-x86_64-netinst.iso

```

---



---

<sup>1</sup>Uninstalling RDO

### 3.2.2.1 Disable firewalld

We need to disable firewalld and replace it with iptables as it is not supported with Havana.<sup>2</sup>

Listing 3.7: Disable firewalld, install iptables

---

```
1 yum install iptables -y
2 systemctl disable firewalld
3 systemctl enable iptables
```

---

### 3.2.2.2 Disable SELinux

The issue with SELinux and OpenStack has been resolved<sup>3</sup>, lets disable it anyway.

Listing 3.8: Disable SELinux

---

```
1 setenforce 0 # does not persist
2 sed -i 's/SELINUX=enforcing/SELINUX=disabled/' /etc/selinux/config
3 # or
4 vi /etc/selinux/config
5 ...
6 SELINUX=disabled
7 ...
```

---

### 3.2.3 SSH Keys

Listing 3.9: SSH Key Copy

---

```
1 echo "StrictHostKeyChecking no" >> /root/.ssh/config
2 for i in {1..3}
3 do
4     ssh-copy-id -i /root/.ssh/id_rsa.pub virospaw01$i
5 done
6 ssh-copy-id -i /root/.ssh/id_rsa.pub virospaw002
```

---

### 3.2.4 Packstack

Listing 3.10: Packstack remote install after big hammer

---

```
1 for i in {0..3}; do ssh virospaw01$i "yum install openstack-packstack
  -y"; done
```

---

---

<sup>2</sup>Red Hat Bugzilla Bug #981583

<sup>3</sup>Red Hat Bugzilla Bug #995780

Listing 3.11: Configure Packstack

---

```
1 DATE=`date +"%Y_%m_%d_%H_%M_%S"`  
2 packstack --gen-answer-file=packstack-${DATE}  
3 vi packstack-${DATE}
```

---

Listing 3.12 is the differences between a newly generated answer file and one that has been configured for GRE. The left side is the original and the right has the changes. Occasionally when the line is too long the change is forced to a new line.

Listing 3.12: Changes in Packstack answer file

---

```
1 CONFIG_CEILOMETER_INSTALL=n  
2 CONFIG_NTP_SERVERS=10.53.252.123,10.53.252.246  
3 CONFIG_NAGIOS_INSTALL=y  
4 CONFIG_GLANCE_HOST=10.53.253.70  
5 CONFIG_CINDER_HOST=10.53.253.70  
6 CONFIG_CINDER_VOLUMES_CREATE=n  
7 CONFIG_NOVA_COMPUTE_HOSTS=10.53.253.70  
8 CONFIG_NEUTRON_SERVER_HOST=10.53.253.100  
9 CONFIG_NEUTRON_L3_HOSTS=10.53.253.100  
10 CONFIG_NEUTRON_DHCP_HOSTS=10.53.253.100  
11 CONFIG_NEUTRON_METADATA_HOSTS=10.53.253.100  
12 CONFIG_NEUTRON_OVS_TENANT_NETWORK_TYPE=gre  
13 CONFIG_NEUTRON_OVS_TUNNEL_RANGES=1:1000  
14 CONFIG_NEUTRON_OVS_TUNNEL_IF=eth1  
15 CONFIG_HORIZON_HOST=10.53.253.120  
16 CONFIG_NAGIOS_HOST=10.53.253.120
```

---

Listing 3.13: Execute Packstack

---

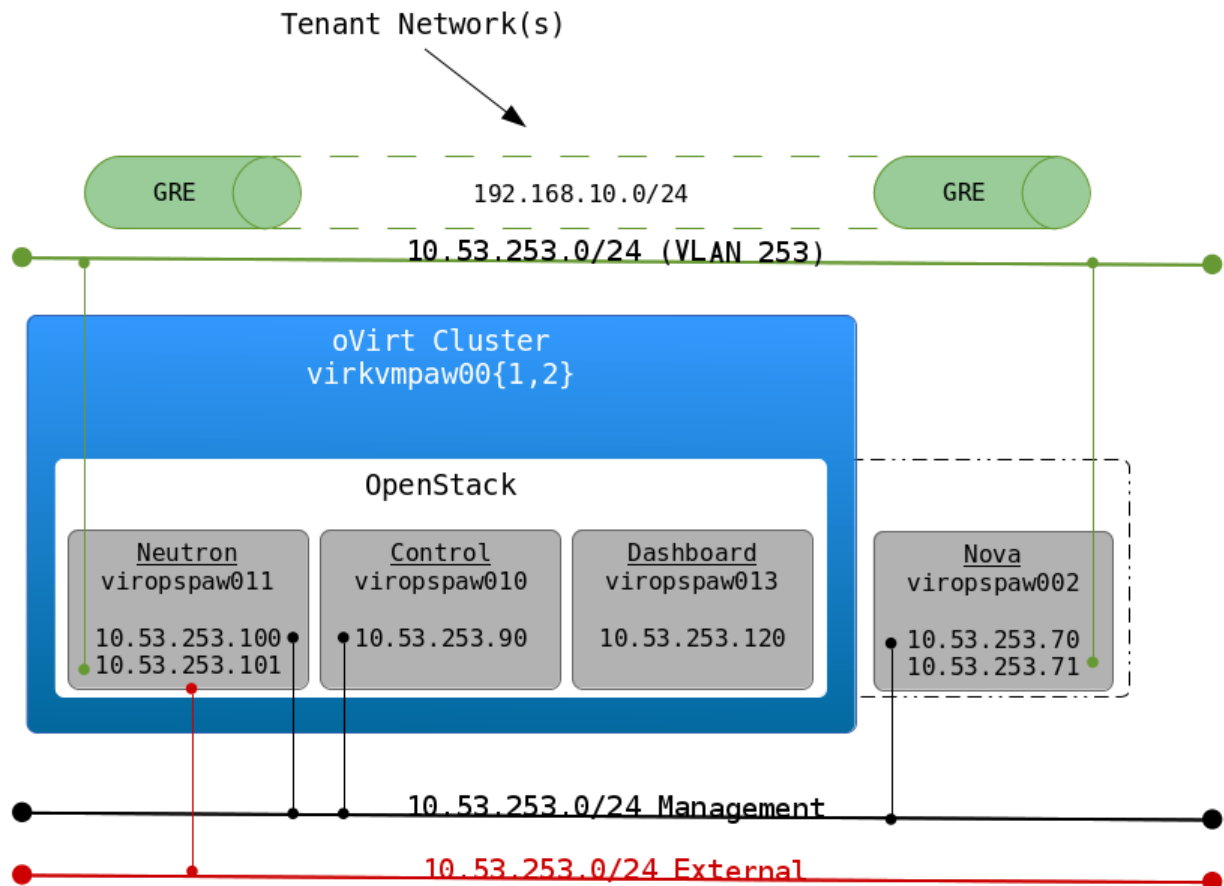
```
1 packstack --answer-file=packstack-${DATE}
```

---

**Note:** Monitor the progress in detail via `/var/tmp/packstack/`

### 3.3 Neutron Networking

Figure 3.1: OpenStack Node Details



Listing 3.14: Nova Node Open vSwitch



---

```

1 ovs-vsctl show
2 95dc7e47-0b2b-49e9-82ae-da4b60dcc68
3     Bridge br-tun
4         Port br-tun
5             Interface br-tun
6                 type: internal
7     Port "gre-1"
8         Interface "gre-1"
9             type: gre
10            options: {in_key=flow, local_ip="10.53.253.71", out_key
11                    =flow, remote_ip="10.53.253.101"}
12     Port patch-int
13         Interface patch-int
14             type: patch
15             options: {peer=patch-tun}
16     Bridge br-int
17         Port patch-tun
18             Interface patch-tun
19                 type: patch
20                 options: {peer=patch-int}
21     Port "qvo7ac76cfd-c8"
22         tag: 4
23         Interface "qvo7ac76cfd-c8"
24     Port "qvoec307994-cd"
25         tag: 4
26         Interface "qvoec307994-cd"
27     Port br-int
28         Interface br-int
29             type: internal
30     ovs_version: "1.11.0"

```

---

Listing 3.15: Neutron Node Open vSwitch

---

```

1 9a0a47e6-e002-4494-b217-5143e9419870
2     Bridge br-int
3         Port "tap5555b73f-79"
4             tag: 1
5             Interface "tap5555b73f-79"
6                 type: internal
7     Port "qr-25c4c000-4b"
8         tag: 1
9         Interface "qr-25c4c000-4b"
10            type: internal
11     Port br-int
12         Interface br-int
13             type: internal
14     Port patch-tun
15         Interface patch-tun
16             type: patch
17             options: {peer=patch-int}

```

---

```

18     Bridge br-tun
19         Port patch-int
20             Interface patch-int
21                 type: patch
22                 options: {peer=patch-tun}
23     Port br-tun
24         Interface br-tun
25             type: internal
26     Port "gre-2"
27         Interface "gre-2"
28             type: gre
29             options: {in_key=flow, local_ip="10.53.253.101",
30                     out_key=flow, remote_ip="10.53.253.71"}
31 Bridge br-ex
32     Port "qg-f7e557b2-b4"
33         Interface "qg-f7e557b2-b4"
34             type: internal
35     Port "eth2"
36         Interface "eth2"
37     Port br-ex
38         Interface br-ex
39             type: internal
40     ovs_version: "1.11.0"

```

---

### 3.3.1 External

Listing 3.16: Create public network and subnet

```

1 neutron net-create public_network --router:external=True
2 neutron subnet-create public_network 10.53.253.0/24 --name
   public_subnet --enable_dhcp=False --allocation-pool start
   =10.53.253.150,end=10.53.253.160 --gateway=10.53.253.1

```

---

### 3.3.2 Tenant

Listing 3.17: Create admin tenant router, network and subnet

```

1 neutron router-create admin_router
2 neutron net-create admin_network
3 neutron subnet-create admin_network 192.168.10.0/24 --name admin_subnet

```

---

Listing 3.18: Set the router gateway

```

1 neutron router-gateway-set admin_router public_network

```

---

### 3.3.3 Troubleshooting

This section is devoted to resolving issues that I encountered while installing and configuring OpenStack.

#### 3.3.3.1 Issue deleting floating IP port

Listing 3.19: Stop Neutron Services

---

```
1 NEUTRON_SERVICES='systemctl --all --full --no-page -t service | grep
  neutron | awk '{print $1}'`
2 for SERVICE in $NEUTRON_SERVICES; do systemctl stop $SERVICE; done
```

---

Listing 3.20: Find the MariaDB root password

---

```
1 grep -i mysql packstack-2013_12_05_10_01_35
2 ...
3 # Username for the MySQL admin user
4 CONFIG_MYSQL_USER=root
5 # Password for the MySQL admin user
6 CONFIG_MYSQL_PW=5729f733fd9f4731
7 ...
```

---

Listing 3.21: Login to MariaDB

---

```
1 mysql -u root -p
```

---

Listing 3.22: Get the ID for the floating IP and delete

---

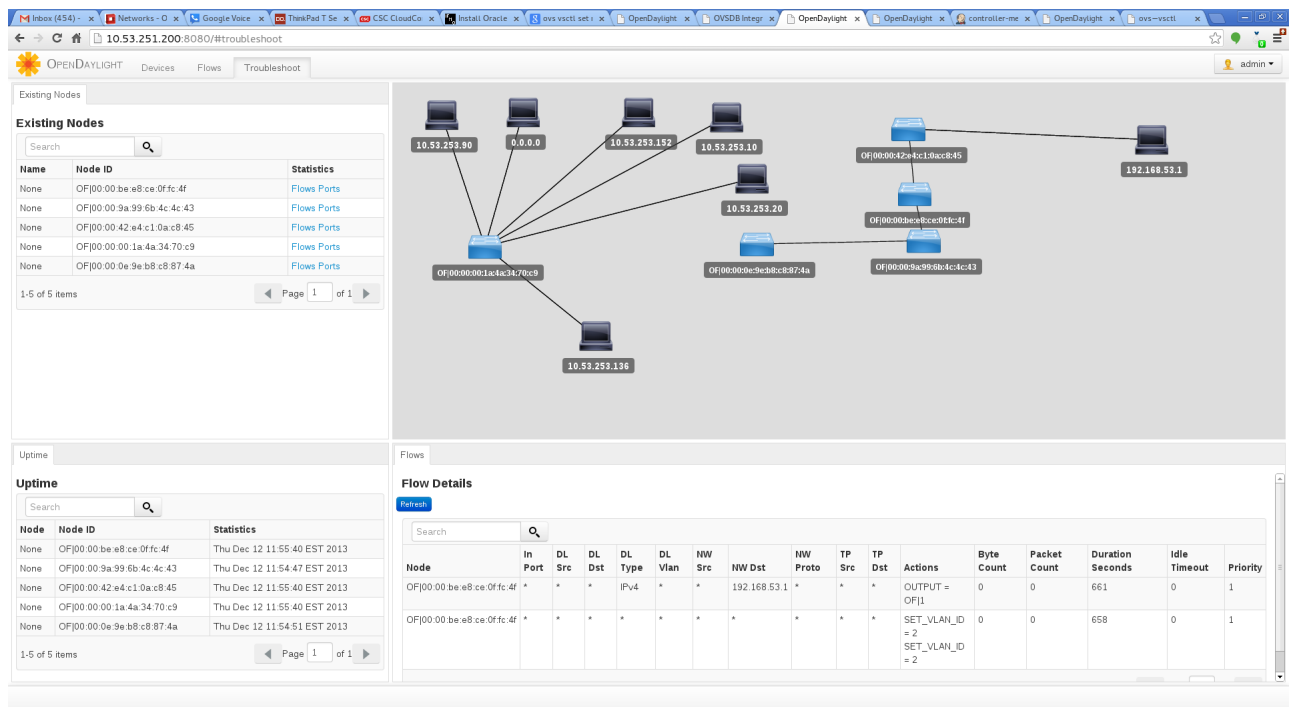
```
1 use ovs_neutron;
2 select * from floatingips;
3 +-----+
4 | id |
5 +-----+
6 | f5af35d0-7a4c-459c-85dc-9f6af9fb1b6c |
7 +-----+
8 delete from floatingips where id = 'f5af35d0-7a4c-459c-85dc-9
  f6af9fb1b6c'
```

---

## 3.4 OpenDaylight

I just started to test with OpenDaylight controller. I am fairly certain that I set this up incorrectly. The only positive is a web-based UI for viewing Open vSwitch / OpenStack networking.

Figure 3.2: OpenDaylight Controller UI



### 3.4.1 Install Controller

I built a new Fedora 19 x64 virtual machine for the OpenDaylight controller.

Listing 3.23: Install Oracle Java JDK

```
1 sudo yum install jdk-7u45-linux-x64.rpm
```

Listing 3.24: Disable firewalld

```
1 systemctl stop firewalld
2 systemctl disable firewalld
```

Listing 3.25: Get the latest controller

```
1 wget https://jenkins.opendaylight.org/controller/job/controller-merge/
   lastSuccessfulBuild/artifact/opendaylight/distribution/opendaylight/
   target/distribution.opendaylight-osgipackage.zip
2 unzip distribution.opendaylight-osgipackage.zip
```

Listing 3.26: Start the controller

```
1 cd opendaylight
2 ./run.sh
```

---

### 3.4.2 Configure Open vSwitch

Listing 3.27: Neutron node set controller and manager

---

```
1 ovs-vsctl set-controller br-tun tcp:10.53.251.200
2 ovs-vsctl set-controller br-ex tcp:10.53.251.200
3 ovs-vsctl set-controller br-int tcp:10.53.251.200
4 ovs-vsctl set-manager tcp:10.53.251.200
```

---

Listing 3.28: Nova node set controller

---

```
1 ovs-vsctl set-controller br-tun tcp:10.53.251.200
2 ovs-vsctl set-controller br-int tcp:10.53.251.200
```

---

# Chapter 4

## Instances

### 4.1 Security Group Configuration

RDO/Packstack creates two "default" security groups. This seems to cause issues with iptables as the instances created are not network available. First we will delete the rules.

Listing 4.1: Delete existing security group rules

---

```
1 SECURITY_GROUP_RULE_LIST='neutron security-group-rule-list \
2 --column id --format csv \
3 | awk 'NR > 1 split($0, a, "\"") {print a[2]}'`
4 for ID in ${SECURITY_GROUP_RULE_LIST};
5 do
6 neutron security-group-rule-delete $ID;
7 done
```

---

Now we will delete the groups themselves. If there is a error stating that the group cannot be deleted because a rule exists that is OK. Just as long as one of the "default" groups has been deleted.

Listing 4.2: Delete existing security groups

---

```
1 SECURITY_GROUP_LIST='neutron security-group-list --column id --format
   csv | awk 'NR > 1 split($0, a, "\"") {print a[2]}'`
2 for ID in ${SECURITY_GROUP_LIST};
3 do
4 neutron security-group-delete $ID;
5 done
```

---

The rules have been deleted lets create our own. Of course these rules open all available ports, probably not what you want to do in production.

---

```
1 neutron security-group-rule-create --protocol icmp --direction ingress
   default
```

---

```

2 neutron security-group-rule-create --protocol tcp --port-range-min 1 --
   port-range-max 65535 --direction ingress default
3 neutron security-group-rule-create --protocol tcp --port-range-min 1 --
   port-range-max 65535 --direction egress default

```

---

## 4.2 Create Images

OpenStack documentation has details for obtaining images for QEMU/KVM Windows and Linux flavors with cloud-init already installed. See link in sources.

### 4.2.1 Ubuntu LTS 12.04

### 4.2.2 Windows 2012 R2

The problem that I ran into was that my / was too small. When using LVM for KVM machines qcow2 images must be converted to raw and copied (dd) into the volume that is created. Glance and/or Cinder use the / filesystem to perform this action (I need to research more information about this). So in order to add Windows I had to manually convert qcow2 image to RAW within a temporary logical volume and then copy into a new logical volume.

Listing 4.3: Windows 2012 R2 Image Create

```

1 gunzip -cd windows_server_2012_r2_standard_eval_kvm_20131031.qcow2.gz |
   glance image-create --name "Windows Server 2012 R2 Std Eval" --disk
   -format=qcow2 --container-format=bare
2 +-----+-----+
3 | Property          | Value                               |
4 +-----+-----+
5 | checksum          | 5fb0a80e56479d31f8b5c1a8e6339206   |
6 | container_format  | bare                                |
7 | created_at        | 2013-11-06T17:35:34                 |
8 | deleted           | False                               |
9 | deleted_at        | None                                |
10 | disk_format        | qcow2                               |
11 | id                 | 32c94fcb-ea3a-4c69-8a49-a6095b8b25ba |
12 | is_public          | False                               |
13 | min_disk           | 0                                    |
14 | min_ram            | 0                                    |
15 | name               | Windows Server 2012 R2 Std Eval     |
16 | owner              | c70e8967165a4b57b8cd0d4265104f01   |
17 | protected          | False                               |
18 | size               | 17182752768                          |
19 | status             | active                              |
20 | updated_at         | 2013-11-06T17:39:48                 |
21 +-----+-----+
22

```

```
23 lvcreate -n temp -L 100G cinder-volumes
24 mkfs.ext3 /dev/cinder-volumes/temp
25 mount /dev/cinder-volumes/temp /mnt/tmp/
26 cd /mnt/tmp/
27 qemu-img convert /var/lib/glance/images/32c94fcb-ea3a-4c69-8a49-
    a6095b8b25ba -O raw ./windows.raw
```

---

## 4.3 Create Instance

Before we create an instance lets import the public key.

Listing 4.4: Add public key

---

```
1 nova keypair-add userkey --pub-key /root/.ssh/id_rsa.pub
```

---

When working with a Windows OS the Administrator password will be set by nova. To retrieve the password use the following command.

Listing 4.5: Windows administrator password

---

```
1 nova get-password windows /root/.ssh/id_rsa
2 38MQ21a1LxYMPH
```

---



# Appendix A

## Notes

### A.1 Neutron

With external egress and ingress virtual machine traffic traversing the Neutron node it should be highly-available. How is this accomplished? What is the length of downtime or in a GRE/VXLAN scenario have multiple "uplinks" to another Neutron host?

---

```
1 Port "gre-1"
2         Interface "gre-1"
3             type: gre
4             options: {in_key=flow, local_ip="10.53.253.71", out_key
                        =flow, remote_ip="10.53.253.101"}
```

---

### A.2 Horizon

How to limit users on Horizon to create instances with new volume.

### A.3 Cinder

When deleting a volume cinder by default will dd /dev/zero this takes forever.

---

Listing A.1: Change Cinder volume clear

---

```
1 sed -i 's/#volume_clear=zero/volume_clear=none/' /etc/cinder/cinder.
   conf
2 # or
3 vi /etc/cinder/cinder.conf
4 ...
5 # Method used to wipe old voumes (valid options are: none,
6 # zero, shred) (string value)
7 volume_clear=none
```

```
8 ...  
9 SERVICES=`systemctl --no-page --full -t service | grep openstack-cinder  
  | awk '{print $1}'`  
10 for CINDER in $SERVICES; do systemctl restart $CINDER; done
```

---

## A.4 Nova

How to properly configure Windows images with the correct QEMU/KVM settings?

# Appendix B

## Links

[Under the hood Open vSwitch](#)

[Neutron with OVS and VLANs](#)

[Install RDO Havana on Fedora 19](#)

[Obtaining Images for OpenStack](#)

[Awk extracting substring using regular expressions with capture groups](#)